

Sreyashi Saha

Matriculation Number : 1747271

Exercise 1 Part A

```
In [1]: import pyspark
        from pyspark.sql import SparkSession
        from pyspark import SparkContext
        from pyspark.sql import SQLContext
```

```
In [2]: list_a = ["spark", "rdd", "python", "context", "create", "class"]
        list_b = ["operation", "apache", "scala", "lambda", "parallel", "partition"]
```

Create two RDD objects of a, b and do the following tasks. Words should be remained in the results of join operations.

```
In [4]: sc = SparkContext()
        a_rdd = sc.parallelize(list_a)
```

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

```
22/07/03 10:33:28 WARN NativeCodeLoader: Unable to load native-hadoop lib
22/07/03 10:33:28 WARN NativeCodeLoader: library for your platform... using builtin-java classes where applicable
22/07/03 10:33:30 WARN Utils: Service 'SparkUI' could not bind on port 40
22/07/03 10:33:30 WARN Utils: 40. Attempting port 4041.
22/07/03 10:33:30 WARN Utils: Service 'SparkUI' could not bind on port 40
22/07/03 10:33:30 WARN Utils: 41. Attempting port 4042.
22/07/03 10:33:30 WARN Utils: Service 'SparkUI' could not bind on port 40
22/07/03 10:33:30 WARN Utils: 42. Attempting port 4043.
22/07/03 10:33:30 WARN Utils: Service 'SparkUI' could not bind on port 40
22/07/03 10:33:30 WARN Utils: 43. Attempting port 4044.
22/07/03 10:33:30 WARN Utils: Service 'SparkUI' could not bind on port 40
22/07/03 10:33:30 WARN Utils: 44. Attempting port 4045.
22/07/03 10:33:30 WARN Utils: Service 'SparkUI' could not bind on port 40
22/07/03 10:33:30 WARN Utils: 45. Attempting port 4046.
22/07/03 10:33:30 WARN Utils: Service 'SparkUI' could not bind on port 40
22/07/03 10:33:30 WARN Utils: 46. Attempting port 4047.
22/07/03 10:33:30 WARN Utils: Service 'SparkUI' could not bind on port 40
22/07/03 10:33:30 WARN Utils: 47. Attempting port 4048.
```

```
In [5]: b_rdd = sc.parallelize(list_b)
```

```
In [7]: a_rdd.collect()
```

```
Out[7]: ['spark', 'rdd', 'python', 'context', 'create', 'class']
```

```
In [8]: b_rdd.collect()
```

```
Out[8]: ['operation', 'apache', 'scala', 'lambda', 'parallel', 'partition']
```

In order to perform the joins we need to have key value pairs. Therefore, we map each word from both the list and assign their values as 1 using the map function.

```
In [9]: # map_a = a_rdd.map(lambda x: (x, 1)).collect()
map_a = a_rdd.map(lambda x: (x, 1))
```

```
In [10]: # map_b = b_rdd.map(lambda x: (x, 1)).collect()
map_b = b_rdd.map(lambda x: (x, 1))
```

```
In [11]: map_a.collect()
```

```
Out[11]: [('spark', 1),
          ('rdd', 1),
          ('python', 1),
          ('context', 1),
          ('create', 1),
          ('class', 1)]
```

```
In [12]: map_b.collect()
```

```
Out[12]: [('operation', 1),
          ('apache', 1),
          ('scala', 1),
          ('lambda', 1),
          ('parallel', 1),
          ('partition', 1)]
```

We can see that all the words in both rdds now have values assigned as 1

1. Perform rightOuterJoin and fullOuterJoin operations between a and b. Briefly explain your solution. (1 point)

```
In [13]: join1 = map_a.rightOuterJoin(map_b)
Right_Outer_Join = join1.map(lambda x: x[0])
```

```
In [14]: Right_Outer_Join.collect()
```

```
Out[14]: ['parallel', 'lambda', 'scala', 'operation', 'apache', 'partition']
```

A right outer join is a method of combining tables. The result includes unmatched rows from only the table that is specified after the RIGHT OUTER JOIN phrase. Here in our case we have map_b after the phrase, it returns all records from the right rdd (map_b), and the matching records from the left rdd (map_a). The result is 0 records from the left side, if there is no match.

```
In [18]: join2 = map_a.fullOuterJoin(map_b)
Full_Outer_join = join2.map(lambda x: x[0]).collect()
```

```
In [16]: Full_Outer_join.collect()
```

```
Out[16]: ['python',
          'spark',
          'context',
          'create',
          'parallel',
          'lambda',
          'class',
          'rdd',
          'scala',
          'operation',
          'apache',
          'partition']
```

The FULL OUTER JOIN keyword returns all records when there is a match in left rdd (map_a) or right rdd(map_b) table records.

```
In [19]: new_list = sc.parallelize(Full_Outer_join)
all_words = new_list.map(lambda word: (word, 1)).collect()
all_words
```

```
Out[19]: [('python', 1),
          ('spark', 1),
          ('context', 1),
          ('create', 1),
          ('parallel', 1),
          ('lambda', 1),
          ('class', 1),
          ('rdd', 1),
          ('scala', 1),
          ('operation', 1),
          ('apache', 1),
          ('partition', 1)]
```

Here I map all the words after performing Full_Outer_join and take them as key value pairs for further implementations.

1. Using map and reduce functions to count how many times the character "s" appears in all a and b. (1 point)

```
In [22]: words = new_list.map(lambda x: (x, 1))
```

In the function below I check if the word received as input has the character 's' in it or not. If it does then I count the number of times I get 's' for that word and then I return the count.

```
In [23]: def check_for_s(temp):
          c=0
          for idx in temp:
              if idx == 's':
                  c+=1
          return c
```

Here I map the words and the number of 's' that the word contains

```
In [24]: count_s_words = words.map(lambda x:(x[0], check_for_s(x[0])))
```

```
In [25]: count_s_words.collect()
```

```
Out[25]: [('python', 0),
          ('spark', 1),
          ('context', 0),
          ('create', 0),
          ('parallel', 0),
          ('lambda', 0),
          ('class', 2),
          ('rdd', 0),
          ('scala', 1),
          ('operation', 0),
          ('apache', 0),
          ('partition', 0)]
```

```
In [27]: count = count_s_words.map(lambda a:a[1]).reduce(lambda a,b:a+b)
```

```
In [29]: print(" The number of times s occurs in the list of words using reduce fu
          The number of times s occurs in the list of words using reduce funtion i
          s 4
```

1. Using aggregate function to count how many times the character "s" appears in all a and b. (1 point)

```
In [30]: count1 = words.aggregate(0,lambda a,x:a+check_for_s(x[0]), lambda a,b: a+
```

```
In [31]: print(" The number of times s occurs in the list of words using aggregate
```

The number of times s occurs in the list of words using aggregate function is 4

In []:

Part b) Basic Operations on DataFrames (6 points)

Use dataset students.json (download from learnweb) for this exercise. First creating DataFrames from the dataset and do several tasks as follows:

```
In [2]: import pyspark
from pyspark.sql import SparkSession
from pyspark import SparkContext
from pyspark.sql import SQLContext
```

```
In [3]: # Read JSON file into dataframe
spark = SparkSession.builder.appName(
    'Read Json File into DataFrame').getOrCreate()
df = spark.read.json("students.json")
df.printSchema()
df.show()
```

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

```
22/07/03 10:58:48 WARN NativeCodeLoader: Unable to load native-hadoop lib
rary for your platform... using builtin-java classes where applicable
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
40. Attempting port 4041.
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
41. Attempting port 4042.
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
42. Attempting port 4043.
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
43. Attempting port 4044.
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
44. Attempting port 4045.
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
45. Attempting port 4046.
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
46. Attempting port 4047.
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
47. Attempting port 4048.
22/07/03 10:58:50 WARN Utils: Service 'SparkUI' could not bind on port 40
48. Attempting port 4049.
```

```
root
```

```
|-- course: string (nullable = true)
|-- dob: string (nullable = true)
|-- first_name: string (nullable = true)
|-- last_name: string (nullable = true)
|-- points: long (nullable = true)
|-- s_id: long (nullable = true)
```

course	dob	first_name	last_name	points	s_id
Humanities and Art	October 14, 1983	Alan	Joe	10	1
Computer Science	September 26, 1980	Martin	Genberg	17	2
Graphic Design	June 12, 1982	Athur	Watson	16	3
Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4
Psychology	November 1, 1978	Kira	Schommer	11	5
Business	17 February 1981	Christian	Kiriam	10	6
Machine Learning	1 January 1984	Barbara	Ballard	14	7
Deep Learning	January 13, 1978	John	null	10	8
Machine Learning	26 December 1989	Marcus	Carson	15	9
Physics	30 December 1987	Marta	Brooks	11	10
Data Analytics	June 12, 1975	Holly	Schwartz	12	11
Computer Science	July 2, 1985	April	Black	null	12
Computer Science	July 22, 1980	Irene	Bradley	13	13
Psychology	7 February 1986	Mark	Weber	12	14
Informatics	May 18, 1987	Rosie	Norman	9	15
Business	August 10, 1984	Martin	Steele	7	16
Machine Learning	16 December 1990	Colin	Martinez	9	17
Data Analytics	null	Bridget	Twain	6	18
Business	7 March 1980	Darlene	Mills	19	19
Data Analytics	June 2, 1985	Zachary	null	10	20

1. Replace the null value(s) in column points by the mean of all points. (0.5 point)

```
In [5]: #Replace mean for null on only points column
mean = df.agg({'points': 'mean'}).collect()
mean = mean[0][0]
print("The mean of the points column is ", mean)
new_df = df.na.fill(value=mean,subset=["points"])
```

The mean of the points column is 11.736842105263158

```
In [6]: new_df.show() # we can see that the null values has been replaced by the
```

course	dob	first_name	last_name	points	s_id
Humanities and Art	October 14, 1983	Alan	Joe	10	1
Computer Science	September 26, 1980	Martin	Genberg	17	2
Graphic Design	June 12, 1982	Athur	Watson	16	3
Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4
Psychology	November 1, 1978	Kira	Schommer	11	5
Business	17 February 1981	Christian	Kiriam	10	6
Machine Learning	1 January 1984	Barbara	Ballard	14	7
Deep Learning	January 13, 1978	John	null	10	8
Machine Learning	26 December 1989	Marcus	Carson	15	9
Physics	30 December 1987	Marta	Brooks	11	10
Data Analytics	June 12, 1975	Holly	Schwartz	12	11
Computer Science	July 2, 1985	April	Black	11	12
Computer Science	July 22, 1980	Irene	Bradley	13	13
Psychology	7 February 1986	Mark	Weber	12	14
Informatics	May 18, 1987	Rosie	Norman	9	15
Business	August 10, 1984	Martin	Steele	7	16
Machine Learning	16 December 1990	Colin	Martinez	9	17
Data Analytics	null	Bridget	Twain	6	18
Business	7 March 1980	Darlene	Mills	19	19
Data Analytics	June 2, 1985	Zachary	null	10	20

1. Replace the null value(s) in column dob and column last name by "unknown" and "--" respectively. (0.5 point)

```
In [8]: df1 = new_df.na.fill(value="unknown",subset=["dob"])
df1.show()
```


course	dob	first_name	last_name	points	s_id
Humanities and Art	October 14, 1983	Alan	Joe	10	1
Computer Science	September 26, 1980	Martin	Genberg	17	2
Graphic Design	June 12, 1982	Athur	Watson	16	3
Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4
Psychology	November 1, 1978	Kira	Schommer	11	5
Business	17 February 1981	Christian	Kiriam	10	6
Machine Learning	1 January 1984	Barbara	Ballard	14	7
Deep Learning	January 13, 1978	John	null	10	8
Machine Learning	26 December 1989	Marcus	Carson	15	9
Physics	30 December 1987	Marta	Brooks	11	10
Data Analytics	June 12, 1975	Holly	Schwartz	12	11
Computer Science	July 2, 1985	April	Black	11	12
Computer Science	July 22, 1980	Irene	Bradley	13	13
Psychology	7 February 1986	Mark	Weber	12	14
Informatics	May 18, 1987	Rosie	Norman	9	15
Business	August 10, 1984	Martin	Steele	7	16
Machine Learning	16 December 1990	Colin	Martinez	9	17
Data Analytics	unknown	Bridget	Twain	6	18
Business	7 March 1980	Darlene	Mills	19	19
Data Analytics	June 2, 1985	Zachary	null	10	20

```
In [9]: df2 = df1.na.fill(value="--",subset=["last_name"])
df2.show()
```

course	dob	first_name	last_name	points	s_id
Humanities and Art	October 14, 1983	Alan	Joe	10	1
Computer Science	September 26, 1980	Martin	Genberg	17	2
Graphic Design	June 12, 1982	Athur	Watson	16	3
Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4
Psychology	November 1, 1978	Kira	Schommer	11	5
Business	17 February 1981	Christian	Kiriam	10	6
Machine Learning	1 January 1984	Barbara	Ballard	14	7
Deep Learning	January 13, 1978	John	--	10	8
Machine Learning	26 December 1989	Marcus	Carson	15	9
Physics	30 December 1987	Marta	Brooks	11	10
Data Analytics	June 12, 1975	Holly	Schwartz	12	11
Computer Science	July 2, 1985	April	Black	11	12
Computer Science	July 22, 1980	Irene	Bradley	13	13
Psychology	7 February 1986	Mark	Weber	12	14
Informatics	May 18, 1987	Rosie	Norman	9	15
Business	August 10, 1984	Martin	Steele	7	16
Machine Learning	16 December 1990	Colin	Martinez	9	17
Data Analytics	unknown	Bridget	Twain	6	18
Business	7 March 1980	Darlene	Mills	19	19
Data Analytics	June 2, 1985	Zachary	--	10	20

1. In the dob column, there exist several formats of dates, e.g. October 14, 1983 and 26 December 1989. Let's convert all the dates into DD-MM-YYYY format where DD, MM and YYYY are two digits for day, two digits for months and four digits for year respectively. (2 points)

```
In [17]: from dateutil import parser
from datetime import datetime
import datetime
```

In this function below I convert the provided date format as "%m-%d-%Y" format and update the values in a new column name new_date. For those rows which have values as 'unknown' I returned "1" as their date of birth.

```
In [21]: def date(input):
    if input == "unknown":
        return "1"
    else:
        obj=str(parser.parse(input))
        d = obj.split(" ")
        d = datetime.datetime.strptime(d[0], '%Y-%m-%d').strftime('%m-%d-%Y')
        return d
```

```
In [22]: from pyspark.sql import functions as F
from pyspark.sql.functions import to_date
from pyspark.sql.types import StringType
```

```
In [23]: data = df2.withColumn("new_date", F.udf(date, StringType())(F.col("dob")))
data.show()
```

new_date	course	dob	first_name	last_name	points	s_id
0-14-1983	Humanities and Art	October 14, 1983	Alan	Joe	10	1
9-26-1980	Computer Science	September 26, 1980	Martin	Genberg	17	2
6-12-1982	Graphic Design	June 12, 1982	Athur	Watson	16	3
4-05-1987	Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4
1-01-1978	Psychology	November 1, 1978	Kira	Schommer	11	5
2-17-1981	Business	17 February 1981	Christian	Kiriam	10	6
1-01-1984	Machine Learning	1 January 1984	Barbara	Ballard	14	7
1-13-1978	Deep Learning	January 13, 1978	John	--	10	8
2-26-1989	Machine Learning	26 December 1989	Marcus	Carson	15	9
2-30-1987	Physics	30 December 1987	Marta	Brooks	11	10
6-12-1975	Data Analytics	June 12, 1975	Holly	Schwartz	12	11
7-02-1985	Computer Science	July 2, 1985	April	Black	11	12
7-22-1980	Computer Science	July 22, 1980	Irene	Bradley	13	13
2-07-1986	Psychology	7 February 1986	Mark	Weber	12	14
5-18-1987	Informatics	May 18, 1987	Rosie	Norman	9	15
8-10-1984	Business	August 10, 1984	Martin	Steele	7	16
2-16-1990	Machine Learning	16 December 1990	Colin	Martinez	9	17
1	Data Analytics	unknown	Bridget	Twain	6	18
3-07-1980	Business	7 March 1980	Darlene	Mills	19	19
6-02-1985	Data Analytics	June 2, 1985	Zachary	--	10	20

1. Insert a new column age and calculate the current age of all students. (1 point)

In [27]: `from dateutil import relativedelta`

In this function I calculate the current age of each student and store them in a new column Current_age. For those rows which have birthdate as "1" I didn't calculate the current age and just returned a null value.

```
In [32]: def calculateAge(birthDate):  
  
    if birthDate == "1":  
#         print('this is 1')  
        return ""  
    else:  
        start_date = datetime.datetime.strptime(birthDate, '%m-%d-%Y')  
        today = datetime.datetime.today()  
        delta = relativedelta.relativedelta(today, start_date)  
        return delta.years
```

```
In [33]: data1 = data.withColumn("Current_age", F.udf(calculateAge, StringType()))  
data1.show()
```

new_date	Current_age	course	dob	first_name	last_name	points	s_id
0-14-1983	38	Humanities and Art	October 14, 1983	Alan	Joe	10	1
9-26-1980	41	Computer Science	September 26, 1980	Martin	Genberg	17	2
6-12-1982	40	Graphic Design	June 12, 1982	Athur	Watson	16	3
4-05-1987	35	Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4
1-01-1978	43	Psychology	November 1, 1978	Kira	Schommer	11	5
2-17-1981	41	Business	17 February 1981	Christian	Kiriam	10	6
1-01-1984	38	Machine Learning	1 January 1984	Barbara	Ballard	14	7
1-13-1978	44	Deep Learning	January 13, 1978	John	--	10	8
2-26-1989	32	Machine Learning	26 December 1989	Marcus	Carson	15	9
2-30-1987	34	Physics	30 December 1987	Marta	Brooks	11	10
6-12-1975	47	Data Analytics	June 12, 1975	Holly	Schwartz	12	11
7-02-1985	37	Computer Science	July 2, 1985	April	Black	11	12
7-22-1980	41	Computer Science	July 22, 1980	Irene	Bradley	13	13
2-07-1986	36	Psychology	7 February 1986	Mark	Weber	12	14
5-18-1987	35	Informatics	May 18, 1987	Rosie	Norman	9	15
8-10-1984	37	Business	August 10, 1984	Martin	Steele	7	16
2-16-1990	31	Machine Learning	16 December 1990	Colin	Martinez	9	17
1		Data Analytics	unknown	Bridget	Twain	6	18
3-07-1980	42	Business	7 March 1980	Darlene	Mills	19	19
6-02-1985	37	Data Analytics	June 2, 1985	Zachary	--	10	20

```
In [35]: sd = data1.agg({'points': 'stddev'}).collect()
print("The standard deviation of the points column is ",sd[0][0])
```

The standard deviation of the points column is 3.246050231475656

```
In [37]: mean = data1.agg({'points': 'mean'}).collect()
print("The mean of the points column is ",mean[0][0])
```

The mean of the points column is 11.7

1. Let's consider granting some points for good performed students in the class. For each student, if his point is larger than 1 standard deviation of all points, then we update his current point to 20, which is the maximum. See Annex 1 for a tutorial on how to calculate standard deviation. (2 points)

```
In [39]: std_1 = mean[0][0]+sd[0][0]
std_2 = mean[0][0]-sd[0][0]
```

Here for each student, I check if his point is larger than 1 standard deviation of all points, then I update his current point to 20.

```
In [41]: from pyspark.sql.functions import when
df3 = data1.withColumn("points", when(data1.points>std_1,20).otherwise(da
df3.show()
```

new_date	Current_age	course	dob	first_name	last_name	points	s_id
0-14-1983	38	Humanities and Art	October 14, 1983	Alan	Joe	10	1 1
9-26-1980	41	Computer Science	September 26, 1980	Martin	Genberg	20	2 0
6-12-1982	40	Graphic Design	June 12, 1982	Athur	Watson	20	3 0
4-05-1987	35	Graphic Design	April 5, 1987	Anabelle	Sanberg	12	4 0
1-01-1978	43	Psychology	November 1, 1978	Kira	Schommer	11	5 1
2-17-1981	41	Business	17 February 1981	Christian	Kiriam	10	6 0
1-01-1984	38	Machine Learning	1 January 1984	Barbara	Ballard	14	7 0
1-13-1978	44	Deep Learning	January 13, 1978	John	--	10	8 0
2-26-1989	32	Machine Learning	26 December 1989	Marcus	Carson	20	9 1
2-30-1987	34	Physics	30 December 1987	Marta	Brooks	11	10 1
6-12-1975	47	Data Analytics	June 12, 1975	Holly	Schwartz	12	11 0
7-02-1985	37	Computer Science	July 2, 1985	April	Black	11	12 0
7-22-1980	41	Computer Science	July 22, 1980	Irene	Bradley	13	13 0
2-07-1986	36	Psychology	7 February 1986	Mark	Weber	12	14 0
5-18-1987	35	Informatics	May 18, 1987	Rosie	Norman	9	15 0
8-10-1984	37	Business	August 10, 1984	Martin	Steele	7	16 0
2-16-1990	31	Machine Learning	16 December 1990	Colin	Martinez	9	17 1
1		Data Analytics	unknown	Bridget	Twain	6	18
3-07-1980	42	Business	7 March 1980	Darlene	Mills	20	19 0
6-02-1985	37	Data Analytics	June 2, 1985	Zachary	--	10	20 0


1. Create a histogram on the new points created in the task 5. (1 point)

```
In [42]: import numpy as np
import matplotlib.pyplot as plt
points=[]
points_array = np.array(df3.select('points').collect())
# points_array
for value in points_array:
    points.append(value[0])
print(points)
student_id=[]
id_array = np.array(df3.select('s_id').collect())
# points_array
for value in id_array:
    student_id.append(value[0])
print(student_id)

plt.hist(points)
# plt.xlabel("sd_id")
plt.ylabel("points")
plt.show()
```

```
[10, 20, 20, 12, 11, 10, 14, 10, 20, 11, 12, 11, 13, 12, 9, 7, 9, 6, 20,
10]
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20]
```



```
In [ ]:
```


Exercise 2

First importing all the libraries required for Spark session in python. A SparkContext represents the connection to a Spark cluster, and is used to create RDDs, accumulators and broadcast variables on that cluster. The sql function on a SQLContext enables applications to run SQL queries programmatically and returns the result as a DataFrame .

```
In [1]: import pyspark
from pyspark import SparkContext
sc = SparkContext()
from pyspark.sql import Row
from pyspark.sql import SQLContext

sqlContext = SQLContext(sc)
from pyspark.sql import SparkSession
```

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

```
22/07/03 08:32:12 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
22/07/03 08:32:13 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.
22/07/03 08:32:13 WARN Utils: Service 'SparkUI' could not bind on port 4041. Attempting port 4042.
22/07/03 08:32:13 WARN Utils: Service 'SparkUI' could not bind on port 4042. Attempting port 4043.
22/07/03 08:32:13 WARN Utils: Service 'SparkUI' could not bind on port 4043. Attempting port 4044.
22/07/03 08:32:13 WARN Utils: Service 'SparkUI' could not bind on port 4044. Attempting port 4045.
22/07/03 08:32:13 WARN Utils: Service 'SparkUI' could not bind on port 4045. Attempting port 4046.
22/07/03 08:32:13 WARN Utils: Service 'SparkUI' could not bind on port 4046. Attempting port 4047.
```

```
/opt/homebrew/lib/python3.9/site-packages/pyspark/sql/context.py:112: FutureWarning: Deprecated in 3.0.0. Use SparkSession.builder.getOrCreate() instead.
  warnings.warn(
```

A SparkSession can be used to create DataFrame, register DataFrame as tables, execute SQL over tables, cache tables, and read data files. Class builder is Builder for SparkSession. appName() sets a name for the application, which will be shown in the Spark web UI. If no application name is set, a randomly generated name will be used. getOrCreate() gets an existing SparkSession or, if there is no existing one, creates a new one based on the options set in this builder.

Using spark.read.csv("path") we can read a CSV file with fields delimited by pipe, comma, tab (and many more) into a Spark DataFrame. These methods take a file path to read from as an argument.

```
In [2]: spark = SparkSession.builder.appName(
        'Read CSV File into DataFrame').getOrCreate()

data = spark.read.csv('ml-10M100K-2/tags.dat', sep='|', inferSchema=True,

# Show the first 5 values of the dataframe
data.show(5)
```

22/07/03 08:32:21 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.

UserID	MovieID	Tag	Timestamp
15	4973	excellent!	1215184630
20	1747	politics	1188263867
20	1747	satire	1188263867
20	2424	chick flick 212	1188263835
20	2424	hanks	1188263835

only showing top 5 rows

```
In [3]: from pyspark import SparkFiles
        from pyspark.sql.functions import col
```

```
In [4]: # Ordering or Sorting the dataframe base din Timestamp column
        data.orderBy(col("Timestamp").asc()).show(10)
```

```

+-----+-----+-----+-----+-----+
|UserID|MovieID|          Tag| Timestamp|
+-----+-----+-----+-----+
| 19106|   1247|    mrs. robinson|1135313387|
| 19106|    145|      car chase|1135313403|
| 19106|   3481|    jack black|1135313411|
| 19106|   3481|       funny|1135313411|
| 19106|  34405|the man they call...|1135313453|
| 19106|   1396|  too many secrets|1135313492|
| 71331|   1396|   setec astronomy|1135361580|
| 22198|   2788|   monty python|1135429210|
| 22198|   1732|   coen brothers|1135429236|
| 22198|   1206|   stanley kubrick|1135429248|
+-----+-----+-----+-----+
only showing top 10 rows

```

```
In [8]: # In order to convert timestamp to readable date we import this in built
        from datetime import datetime
```

In this function I take my timestamp value from the dataframe which is in string and convert it into integer. After that I use `datetime.utcfromtimestamp().strftime()` function to convert the time stamp into a format as day-month-Year Hours:Minutes:Seconds. Then I return this readable date.

```
In [9]: def date_time(timestamp):
        timesatmp = int (timestamp)
        date = datetime.utcfromtimestamp(timesatmp).strftime('%d-%m-%Y %H:%M:
        return date
```

```
In [10]: date_time("1135429236")
```

```
Out[10]: '24-12-2005 13:00:36'
```

```
In [11]: from pyspark.sql import functions as F
        from pyspark.sql.functions import to_date
        from pyspark.sql.types import StringType
```

In the code below I have taken the Timestamp column and performed the timestamp to date conversion by calling my user defined program `date_time` and saved the returned value from the function in a separate column named as `Date_Time`. We can see that we have obtained a new column which contains the new readable date and time.

```
In [12]: data1 = data.withColumn("Date_Time", F.udf(date_time, StringType())(F.co
        data1.show())
```

```
[Stage 3:>                                (0 +
1) / 1]
```

UserID	MovieID	Tag	Timestamp	Date_Time
15	4973	excellent!	1215184630	04-07-2008 15:17:10
20	1747	politics	1188263867	28-08-2007 01:17:47
20	1747	satire	1188263867	28-08-2007 01:17:47
20	2424	chick flick 212	1188263835	28-08-2007 01:17:15
20	2424	hanks	1188263835	28-08-2007 01:17:15
20	2424	ryan	1188263835	28-08-2007 01:17:15
20	2947	action	1188263755	28-08-2007 01:15:55
20	2947	bond	1188263756	28-08-2007 01:15:56
20	3033	spoof	1188263880	28-08-2007 01:18:00
20	3033	star wars	1188263880	28-08-2007 01:18:00
20	7438	bloody	1188263801	28-08-2007 01:16:41
20	7438	kung fu	1188263801	28-08-2007 01:16:41
20	7438	Tarantino	1188263801	28-08-2007 01:16:41
21	55247	R	1205081506	09-03-2008 16:51:46
21	55253	NC-17	1205081488	09-03-2008 16:51:28
25	50	Kevin Spacey	1166101426	14-12-2006 13:03:46
25	6709	Johnny Depp	1162147221	29-10-2006 18:40:21
31	65	buddy comedy	1188263759	28-08-2007 01:15:59
31	546	strangely compelling	1188263674	28-08-2007 01:14:34
31	1091	catastrophe	1188263741	28-08-2007 01:15:41

only showing top 20 rows

1. A tagging session for a user can be defined as the duration in which he/she generated tagging activities. Typically, an inactive duration of 30 mins is considered as a termination of the tagging session. Your task is to separate out tagging sessions for each user.

```
In [14]: from pyspark.sql import Window
```

PYSPARK LAG is a function in PySpark that works as the offset row returning the value of the before row of a column with respect to the current row in PySpark. This is an operation in PySpark that returns the row just before the current row. Here I have used lag in the timestamp column and partitioned by the userIds. Here I get the previous timestamp for each userid.

```
In [28]: data1 = data1.withColumn(
    "session_id",
    F.lag("Timestamp", 1).over(Window.partitionBy("UserID").orderBy("Times
    )
```

```
In [29]: data1.show(10)
```

```
[Stage 37:> (0 + 1) / 1]
```

UserID	MovieID	Tag	Timestamp	Date_Time	session_id
20	2947	action	1188263755	28-08-2007 01:15:55	null
20	2947	bond	1188263756	28-08-2007 01:15:56	1188263755
20	7438	bloody	1188263801	28-08-2007 01:16:41	1188263756
20	7438	kung fu	1188263801	28-08-2007 01:16:41	1188263801
20	7438	Tarantino	1188263801	28-08-2007 01:16:41	1188263801
20	2424	chick flick 212	1188263835	28-08-2007 01:17:15	1188263801
20	2424	hanks	1188263835	28-08-2007 01:17:15	1188263835
20	2424	ryan	1188263835	28-08-2007 01:17:15	1188263835
20	1747	politics	1188263867	28-08-2007 01:17:47	1188263835
20	1747	satire	1188263867	28-08-2007 01:17:47	1188263867

only showing top 10 rows

Define if the session is the 1st one (more than 1800s after the previous one). Here I subtract the timestamps from the timestamp column and session id column created and check if it is less than 1800 seconds or 30 minutes. If it is less then I assign it as 0 or else I assign it as 1.

```
In [30]: df = data1.withColumn("session_id",
    F.when(F.col("Timestamp") - F.col("session_id") <= 1800, 0).otherwise
    )
```

```
In [31]: df.show(10)
```

```
[Stage 40:> (0 + 1) / 1]
```

```

+-----+-----+-----+-----+-----+-----+
+
|UserID|MovieID|          Tag| Timestamp|          Date_Time|session_id
|
+-----+-----+-----+-----+-----+-----+
+
|    20|    2947|      action|1188263755|28-08-2007 01:15:55|      1
|
|    20|    2947|      bond|1188263756|28-08-2007 01:15:56|      0
|
|    20|    7438|    bloody|1188263801|28-08-2007 01:16:41|      0
|
|    20|    7438|    kung fu|1188263801|28-08-2007 01:16:41|      0
|
|    20|    7438|   Tarantino|1188263801|28-08-2007 01:16:41|      0
|
|    20|    2424|chick flick 212|1188263835|28-08-2007 01:17:15|      0
|
|    20|    2424|      hanks|1188263835|28-08-2007 01:17:15|      0
|
|    20|    2424|      ryan|1188263835|28-08-2007 01:17:15|      0
|
|    20|    1747|    politics|1188263867|28-08-2007 01:17:47|      0
|
|    20|    1747|      satire|1188263867|28-08-2007 01:17:47|      0
|
+-----+-----+-----+-----+-----+-----+
+
only showing top 10 rows

```

Here I find out the unique session id based on the above result. So basically I sum over the session Id column and partition by User Id and again order by the timestamp column. Note that we can get same ids for different users.

```

In [32]: # create a unique id per session
df = df.withColumn(
    "session_id",
    F.sum("session_id").over(Window.partitionBy("userid").orderBy("timest
    ))

```

```

In [33]: df.show(10)

```

```

[Stage 43:>                                (0 +
1) / 1]

```

```

+-----+-----+-----+-----+-----+-----+
+
|UserID|MovieID|          Tag| Timestamp|          Date_Time|session_id
|
+-----+-----+-----+-----+-----+-----+
+
|    20|    2947|      action|1188263755|28-08-2007 01:15:55|      1
|
|    20|    2947|      bond|1188263756|28-08-2007 01:15:56|      1
|
|    20|    7438|    bloody|1188263801|28-08-2007 01:16:41|      1
|
|    20|    7438|    kung fu|1188263801|28-08-2007 01:16:41|      1
|
|    20|    7438|    Tarantino|1188263801|28-08-2007 01:16:41|      1
|
|    20|    2424|chick flick 212|1188263835|28-08-2007 01:17:15|      1
|
|    20|    2424|      hanks|1188263835|28-08-2007 01:17:15|      1
|
|    20|    2424|      ryan|1188263835|28-08-2007 01:17:15|      1
|
|    20|    1747|    politics|1188263867|28-08-2007 01:17:47|      1
|
|    20|    1747|      satire|1188263867|28-08-2007 01:17:47|      1
|
+-----+-----+-----+-----+-----+-----+
+
only showing top 10 rows

```

In the code below I finally create a unique id per session per user where F.dense_rank() returns the rank of rows within a window partition without any gaps.

```
In [34]: df = df.withColumn(
          "session_id", F.dense_rank().over(Window.orderBy("userid", "session_id"))
        )
```

```
In [36]: df.show(10)
```

```

22/07/03 10:19:04 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:19:04 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:19:04 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

```

```
[Stage 47:>                                     (0 + 1) / 1]
```

22/07/03 10:19:06 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:06 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:06 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

[Stage 47:> (0 + 1) / 1]

22/07/03 10:19:13 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:13 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:13 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

```
+-----+-----+-----+-----+-----+-----+
+
|UserID|MovieID|          Tag| Timestamp|          Date_Time|session_id
|
+-----+-----+-----+-----+-----+-----+
+
|    15|    4973|    excellent!|1215184630|04-07-2008 15:17:10|          1
|
|    20|    2947|         action|1188263755|28-08-2007 01:15:55|          2
|
|    20|    2947|         bond|1188263756|28-08-2007 01:15:56|          2
|
|    20|    7438|        bloody|1188263801|28-08-2007 01:16:41|          2
|
|    20|    7438|       kung fu|1188263801|28-08-2007 01:16:41|          2
|
|    20|    7438|    Tarantino|1188263801|28-08-2007 01:16:41|          2
|
|    20|    2424|chick flick 212|1188263835|28-08-2007 01:17:15|          2
|
|    20|    2424|         hanks|1188263835|28-08-2007 01:17:15|          2
|
|    20|    2424|         ryan|1188263835|28-08-2007 01:17:15|          2
|
|    20|    1747|       politics|1188263867|28-08-2007 01:17:47|          2
|
+-----+-----+-----+-----+-----+-----+
+
only showing top 10 rows
```

1. Once you have all the tagging sessions for each user, calculate the frequency of tagging for each user session.


```
In [37]: # Here I have obtained the frequency of tagging for each user session
df1 = df.groupby(['UserID', 'session_id']).count()
df1.show()
```

22/07/03 10:19:19 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:19 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:19 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:19 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:20 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:19:20 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

UserID	session_id	count
15	1	1
20	2	12
21	3	2
25	4	1
25	5	1
31	6	5
32	7	1
39	8	5
48	9	2
49	10	15
75	11	1
78	12	1
109	13	11
109	14	2
109	15	3
109	16	1
109	17	1
109	18	1
109	19	4
109	20	1

only showing top 20 rows

1. Find a mean and standard deviation of the tagging frequency of each user.

```
In [40]: df2 = df1.groupby('UserID', 'session_id').agg({'count': 'mean'})
df2.show()
```

22/07/03 10:20:43 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:20:43 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:20:43 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:20:43 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:20:44 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:20:44 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

```
+-----+-----+-----+
|UserID|session_id|avg(count)|
+-----+-----+-----+
|    15|         1|        1.0|
|    20|         2|       12.0|
|    21|         3|        2.0|
|    25|         4|        1.0|
|    25|         5|        1.0|
|    31|         6|        5.0|
|    32|         7|        1.0|
|    39|         8|        5.0|
|    48|         9|        2.0|
|    49|        10|       15.0|
|    75|        11|        1.0|
|    78|        12|        1.0|
|   109|        13|       11.0|
|   109|        14|        2.0|
|   109|        15|        3.0|
|   109|        16|        1.0|
|   109|        17|        1.0|
|   109|        18|        1.0|
|   109|        19|        4.0|
|   109|        20|        1.0|
+-----+-----+-----+
```

only showing top 20 rows

```
In [41]: df3 = df1.groupby('UserID').agg({'session_id': 'stddev'}).show()
```

22/07/03 10:21:03 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:21:03 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:21:04 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:21:04 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:21:04 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:21:04 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

```
+-----+-----+
|UserID|stddev(session_id)|
+-----+-----+
|    15|                null|
|    20|                null|
|    21|                null|
|    25|0.7071067811865476|
|    31|                null|
|    32|                null|
|    39|                null|
|    48|                null|
|    49|                null|
|    75|                null|
|    78|                null|
|   109|2.7386127875258306|
|   127|                null|
|   133|                null|
|   146| 96.27304918823336|
|   147|                null|
|   170|                null|
|   175|0.7071067811865476|
|   181|                null|
|   190|1.2909944487358056|
+-----+-----+
```

only showing top 20 rows

1. Find a mean and standard deviation of the tagging frequency for across users.

```
In [42]: mean = df1.agg({'count': 'mean'}).collect()
print("The mean across all users = ",mean[0][0])
```

```

22/07/03 10:21:29 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:29 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:29 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:29 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:30 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:30 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
The mean across all users = 7.300084014358817

```

```

In [43]: stddev = df1.agg({'count': 'stddev'}).collect()
print("The standard deviation across all users = ",stddev[0][0])

```

```

22/07/03 10:21:32 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:32 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:32 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:32 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:33 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
22/07/03 10:21:33 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.
The standard deviation across all users = 22.264293050264985

```

1. Provide the list of users with a mean tagging frequency within the two standard deviation from the mean frequency of all users.

```

In [45]: # Range of two standard deviation from mean frequency of all users
std_1 = mean[0][0]+2*stddev[0][0]
std_2 = mean[0][0]-2*stddev[0][0]

```

Here I check if the tagging frequency lies between the above range of values. If it does then I assign it as 1 or else I assign it as 0

```
In [46]: df1 = df1.withColumnRenamed("count", "Frequency")
from pyspark.sql.functions import when
df4 = df1.withColumn("count", when(df1.Frequency < std_1, when(df1.Freque
df4.show(100)
```

22/07/03 10:26:01 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:26:01 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:26:01 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:26:01 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:26:02 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:26:02 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

UserID	session_id	Frequency	count
15	1	1	1
20	2	12	1
21	3	2	1
25	4	1	1
25	5	1	1
31	6	5	1
32	7	1	1
39	8	5	1
48	9	2	1
49	10	15	1
75	11	1	1
78	12	1	1
109	13	11	1
109	14	2	1
109	15	3	1
109	16	1	1
109	17	1	1
109	18	1	1
109	19	4	1
109	20	1	1
109	21	1	1
127	22	26	1
133	23	5	1
146	24	1	1
146	25	1	1
146	26	2	1
146	27	12	1
146	28	2	1
146	29	4	1
146	30	18	1

146	31	53	0
146	32	3	1
146	33	3	1
146	34	2	1
146	35	2	1
146	36	2	1
146	37	1	1
146	38	1	1
146	39	1	1
146	40	4	1
146	41	1	1
146	42	1	1
146	43	2	1
146	44	4	1
146	45	2	1
146	46	2	1
146	47	2	1
146	48	8	1
146	49	4	1
146	50	1	1
146	51	1	1
146	52	1	1
146	53	1	1
146	54	10	1
146	55	1	1
146	56	1	1
146	57	4	1
146	58	1	1
146	59	18	1
146	60	1	1
146	61	6	1
146	62	4	1
146	63	1	1
146	64	1	1
146	65	3	1
146	66	14	1
146	67	69	0
146	68	3	1
146	69	39	1
146	70	27	1
146	71	7	1
146	72	4	1
146	73	2	1
146	74	4	1
146	75	8	1
146	76	1	1
146	77	40	1
146	78	2	1
146	79	3	1
146	80	1	1
146	81	3	1
146	82	1	1
146	83	63	0
146	84	56	0
146	85	5	1
146	86	4	1
146	87	2	1

	146		88		2		1	
	146		89		1		1	
	146		90		2		1	
	146		91		2		1	
	146		92		3		1	
	146		93		1		1	
	146		94		4		1	
	146		95		2		1	
	146		96		2		1	
	146		97		1		1	
	146		98		3		1	
	146		99		6		1	
	146		100		16		1	

+-----+-----+-----+-----+

only showing top 100 rows

```
In [47]: df4 = df4.withColumnRenamed("count", "Users_in_the_range")
df5 = df4.filter(df4.Users_in_the_range == 1)
```

this dataframe df5 contains all the users whose mean frequency lies within 2 standard devaitions away from mean frequency of all users.

```
In [48]: df5.show(10)
```

22/07/03 10:27:19 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:27:19 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:27:19 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:27:19 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:27:20 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

22/07/03 10:27:20 WARN WindowExec: No Partition Defined for Window operation! Moving all data to a single partition, this can cause serious performance degradation.

UserID	session_id	Frequency	Users_in_the_range
15	1	1	1
20	2	12	1
21	3	2	1
25	4	1	1
25	5	1	1
31	6	5	1
32	7	1	1
39	8	5	1
48	9	2	1
49	10	15	1

only showing top 10 rows

In []: